

MNIST | Cross-modal generation
Five fixed held-out pairs · both directions · numbers: per-image SSIM
* DiT normalization recovered from training log (3 decimals).

